

Figure 1: LoRA Fine-tuning Convergence (Qwen2.5-7B, 872 examples)

